

Homework 2

Due 2/6 12pm

1.

$$P(r|\lambda) = \exp(-\lambda) \frac{\lambda^r}{r!}$$

(a) Taking the log of the likelihood function:

$$L(\lambda) = \ln P(r|\lambda) = -\lambda + r \ln \lambda - \ln r!$$

finding the maximum by taking the derivative with respect to λ and setting it to zero:

$$\frac{dL}{d\lambda} = -1 + \frac{r}{\lambda} = 0 \implies \hat{\lambda} = r$$

so the maximum likelihood estimate for λ is $\hat{\lambda} = r$.

(b) Given the derivative with respect to the function $\ln \lambda$:

$$\frac{d}{d(\ln \lambda)} u^n = n u^n, \quad \frac{d}{d(\ln \lambda)} \ln \lambda = 1$$

we can find the curvature of the log likelihood function:

$$\begin{aligned} \frac{d}{d(\ln \lambda)} L(\lambda) &= -\lambda + r = 0 \implies \hat{\lambda} = r \\ \frac{d^2}{d(\ln \lambda)^2} L(\lambda) &= -\lambda = k \end{aligned}$$

For a normal distribution with width σ , the curvature is $k = -1/\sigma^2$. So the width is approximately

$$\sigma \propto \frac{1}{\sqrt{-k}} = \frac{1}{\sqrt{\lambda}}$$

and the 95% confidence interval at the MLE is approximately

$$\hat{\lambda} \pm 2\sigma = r \pm \frac{2}{\sqrt{\hat{\lambda}}}$$

(c) Given the new Poisson distribution

$$P(r|\lambda) = \exp(-(\lambda + b)) \frac{(\lambda + b)^r}{r!}$$

the log likelihood function is

$$L(\lambda) = -(\lambda + b) + r \ln(\lambda + b) - \ln r!$$

and the maximum likelihood estimate for λ is

$$\frac{dL}{d\lambda} = -1 + \frac{r}{\lambda + b} = 0 \implies \hat{\lambda} = r - b$$

the value $\hat{\lambda} = 9 - 13 = -4$ is not physically meaningful, so the MLE will be the lowest possible value for λ which is $\hat{\lambda} = 0$. From this we can infer that the remote star is very dim. The Bayesian posterior distribution for λ is

$$P(\lambda|r) = \frac{P(r|\lambda)P(\lambda)}{P(r)}$$

and sketched in the figure below




Figure 1.1: The posterior distribution for λ given r




Figure 1.2: Segment of Gaussian distribution at $\{x_1, y_1\}$

2. (a) From the geometric picture as shown in Figure 1.2, the segment of the Gaussian is a circle with radius $\rho = \sqrt{x_1^2 + y_1^2}$ and the circumference is $2\pi\rho$ which directly relates to the extra factor of ρ and canceling the 2π in the denominator. This is also related to the Jacobian when transforming from Cartesian to polar coordinates when computing the integral. Using the integral of a 1D Gaussian:

$$\int_{-\infty}^{\infty} \exp(-ax^2) dx = \sqrt{\frac{\pi}{a}}$$

Verifying that the integral is normalized:

$$\begin{aligned} \frac{1}{2\pi\sigma^2} \iint_{-\infty}^{\infty} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right) dx dy &= \frac{1}{\sigma^2} \int_0^{\infty} \exp\left(-\frac{\rho^2}{2\sigma^2}\right) \rho d\rho \\ \frac{1}{2\pi\sigma^2} \int_{-\infty}^{\infty} \exp\left(-\frac{x^2}{2\sigma^2}\right) dx \int_{-\infty}^{\infty} \exp\left(-\frac{y^2}{2\sigma^2}\right) dy &= \frac{1}{\sigma^2} \int_0^{\infty} \exp\left(-\frac{\rho^2}{2\sigma^2}\right) \rho d\rho \\ \frac{1}{2\pi\sigma^2} \sqrt{2\pi\sigma^2} \sqrt{2\pi\sigma^2} &= \frac{1}{\sigma^2} \int_0^{\infty} \sigma^2 \exp(-u) du \\ \frac{2\pi\sigma^2}{2\pi\sigma^2} &= \left[-e^{-u}\right]_0^{\infty} = 1 \end{aligned}$$

so both integrals are normalized.

(b)

$$P(\rho) = \frac{\rho}{\sigma_w^k} \exp\left(-\frac{\rho^2}{2\sigma_w^2}\right)$$

The sketch for the distribution of $P(\rho)$ is shown in Figure 1.3.




Figure 1.3: The distribution of ρ for $k = 1000$

(c) The standard deviation is

$$\sigma_w = \sqrt{\frac{\sum (w - \mu)^2}{k}}$$

and because the distribution is centered at the origin, the mean is $\mu = 0$, so

$$\sigma_w = \sqrt{\frac{\rho^2}{k}}$$

Since most of the probability mass lies around ρ or equivalently a radius of the thin shell where $\rho = r = \sigma_w \sqrt{k}$, and the thickness of the shell is equivalent to the standard deviation of the gaussian $= r/\sqrt{k}$.

(d) Taking the ratio of the probability density from the origin to a point $\rho = \omega_w \sqrt{k}$ away:

$$\frac{P(0)}{P(\sigma_w \sqrt{k})} = \exp\left\{\frac{\sigma_w^2 k}{2\sigma_w^2}\right\} = \exp\left\{\frac{k}{2}\right\}$$

(e) For a shell to contain 95% of the probability mass, $\sigma_w = 2$ and thus the radius and thickness of the shell are

$$r = \sigma_w \sqrt{k} = 2\sqrt{1000} = 63.25, \quad \frac{r}{\sqrt{1000}} = 2$$

and the probability density is $\exp\{500\} \approx 10^{217}$ larger at the origin than at the edge of the shell.

(f) For a 1% difference in σ_w at the origin, the radius term is at zero so the exponents are 1, so the ratio of the probability densities is

$$\frac{(1.01\sigma_w)^k}{\sigma_w^k} = 1.01^k \approx 20959.$$

(g) Because of the large amount of parameters, the MLE would have an expected value at where the probability density is the largest, which is at the origin, but as we have seen, the probability mass is almost all at the edge of the thin shell of radius $\sigma_w \sqrt{k}$. The narrow peak at the origin will overwhelm the MLE and thus we don't see the full picture of the distribution.

3. (a) From class

$$\mathcal{R} = \frac{\frac{F_a!F_b!}{(F+1)!}}{1/2^F} = \frac{2^F F_a!F_b!}{(F+1)!}$$

(b) Taking the log of the ratio and using Stirling's approximation as shown in the code we get the plot of the three simulations in Figure 1.4




Figure 1.4: The log of the ratio of the likelihoods for the three simulations

(c) The trajectory of the two biased coin models are shown in Figure 1.5 and 1.6.

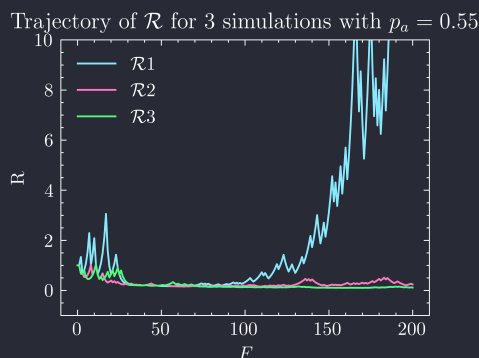


Figure 1.5: There isn't clear evidence for a bias in the first 200 flips




Figure 1.6: A clear bias is shown in the first 200 flips here

and the posterior distribution for the two biased coin models are shown in Figure 1.7 and 1.8.

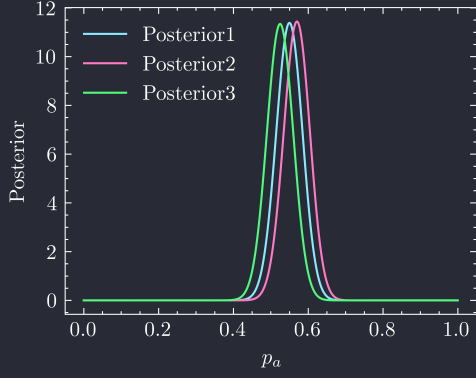


Figure 1.7: We can see that the distribution is centered around roughly $p_a = 0.55$ as expected

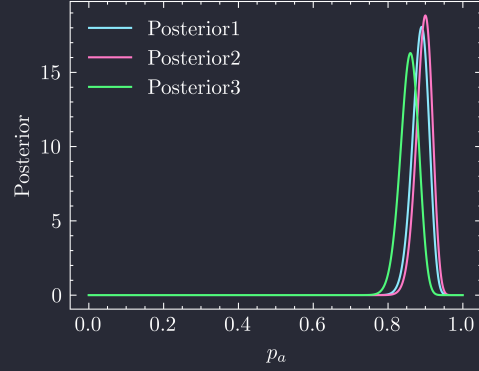


Figure 1.8: There is definitely a bias towards $p_a = 0.9$ in this distribution

(d) From the Gaussian, we can approximate the error bars as $p_a - \frac{1}{2}$, so

$$p_a - 0.5 = \frac{\sigma}{\sqrt{F}} \rightarrow F = \frac{\sigma^2}{(p_a - 0.5)^2}$$

so we can get a rough estimate of the number of flips needed to distinguish between the two models for within 2-3 standard deviations as shown in 1.9 and 1.10.

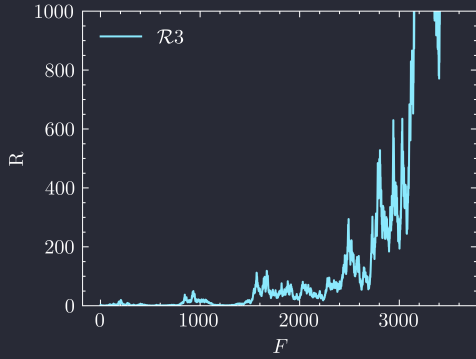


Figure 1.9: For $\sigma = 3$ the model is clearly distinguishable after ≈ 3600 flips




Figure 1.10: For $\sigma = 2$ the model does seem to be biased, but it is not as confident as the previous case

(e) It is surprising to find out that finding evidence against \mathcal{H}_1 would be slower than finding evidence for it, but in hindsight it makes sense for cases where the biased probabilities are close to the unbiased coin model since we would need an exponential number of flips as the bent coin model probabilities is closer to $p_o = 1/2$.

From the log of the ratio

$$\begin{aligned} \ln \mathcal{R} &= F \ln 2 + \ln F_a! + \ln F_b! - \ln(F+1)! \\ &= F \ln 2 + \ln(F - F_b)! + \ln(F - F_a)! - \ln(F+1)! \\ &= F \ln 2 + (F - F_b) \ln(F - F_b) - (F - F_b) + \frac{1}{2} \ln 2\pi(F - F_b) \dots \end{aligned}$$

taking the derivative

$$\begin{aligned} \frac{d}{dF} \ln \mathcal{R} &= \ln 2 + 1 + \ln(F - F_b) - 1 + \frac{1}{2} \frac{1}{F - F_b} \dots \\ &= \ln 2 + \ln(F_a) + \frac{1}{2F_a} + \ln(F_b) + \frac{1}{2F_b} - \ln(F+1) - \frac{1}{2(F+1)} \end{aligned}$$

The fastest $\log \mathcal{R}$ can grow is when F is very large, and the fractional terms at very large F will be negligible and the log terms grow much slower (e.g. $\ln(10^{20}) \approx 46$), so the derivative will be approximately a constant which relates to a linear growth. The fastest $\log \mathcal{R}$ can fall is when F is small and where the $-\ln(F+1)$ term will dominate thus the growth will be approximately logarithmic.

PYTHON CODE BELOW

```
1 # %%
2 import numpy as np
3 import scipy as sp
4 import matplotlib.pyplot as plt
5 import matplotlib as mpl
6 import scienceplots
7 import gmpy2 as gm
8
9 # Science plot package + Dracula theme
10 plt.style.use(['science', 'dark_background'])
11 plt.rcParams['axes.facecolor'] = '#282a36'
12 plt.rcParams['figure.facecolor'] = '#282a36'
13 colorcycle = ['#8be9fd', '#ff79c6', '#50fa7b', '#bd93f9', '#ffb86c', '#ff5555', '#f1fa8c',
14              '#6272a4']
15 plt.rcParams['axes.prop_cycle'] = mpl.cycler(color=colorcycle)
16 white = '#f8f8f2' # foreground
17
18 # change dpi
19 plt.rcParams['figure.dpi'] = 1024
20
21 # %%
22 # 3b
23 # Calculating the ratio R
24 # Stirling's approximation
25 def power(n, exp):
26     result = 1
27     for _ in range(exp):
28         result = gm.mul(n, result)
29     return result
30
31 def fact(n):
32     if n < 10:
33         return np.math.factorial(n)
34     a = power(n, n)
35     b = gm.exp(-n)
36     c = gm.sqrt(2 * np.pi * n)
37     return a * b * c
38
39 def logratio(F_a, F_b):
40     return (F_a + F_b) * gm.log(2) + gm.log(fact(F_a)) + gm.log(fact(F_b)) - gm.log(fact(F_a + F_b + 1))
41
42 def flip():
43     if np.random.rand() < 0.5:
44         return 1
45     else:
46         return 0
47
48 def simulate(n):
49     F_a = 0
50     F_b = 0
51     R = [1]
52     for i in range(n):
53         if flip():
54             F_a += 1
55         else:
56             F_b += 1
57         R.append(logratio(F_a, F_b))
58     return R
59 print(fact(100))
60 print(2 * gm.log(fact(100)))
61 print(gm.log(fact(201)))
62 print(logratio(100, 100))
63 # numpy array with 1 to 200 on the x-axis
64 n = np.arange(0, 201)
65 # calculate the ratio for each n
66 # R = simulate(200)
```

```

67 # simulating 3 times and plotting them side by side
68 plt.figure()
69 for i in range(3):
70     R = simulate(200)
71     # Plot on the first subplot
72     plt.plot(n, R, label='log  $\mathcal{R}$ ' + str(i+1))
73     plt.xlabel('$F$')
74     plt.ylabel('R')
75 plt.title('Part 3b: log  $\mathcal{R}$  for 3 simulations')
76 plt.legend()
77 plt.show()
78
79 # plot the ratio
80 # plt.plot(n, R, label='R')
81 # plt.xlabel('$F$')
82 # plt.ylabel('R')
83 # plt.legend()
84 # plt.show()
85
86 # %%
87 # for p_a = 0.55
88 def bflip():
89     if np.random.rand() < 0.55:
90         return 1
91     else:
92         return 0
93
94 def vbflip():
95     if np.random.rand() < 0.9:
96         return 1
97     else:
98         return 0
99
100 Fa_toss_count_55 = []
101 Fa_toss_count_90 = []
102
103 def ratio(F_a, F_b):
104     return fact(F_a) * fact(F_b) / fact(F_a + F_b + 1) * 2 ** (F_a + F_b)
105
106 def simulate_b(n):
107     F_a = 0
108     F_b = 0
109     R = [1]
110     for i in range(n):
111         if bflip():
112             F_a += 1
113         else:
114             F_b += 1
115         R = np.append(R, ratio(F_a, F_b))
116     return R, F_a
117
118 print(2 ** 100 * gm.factorial(90) * gm.factorial(10) / gm.factorial(101))
119
120 print(2 ** 500 * gm.factorial(275) * gm.factorial(225) / gm.factorial(501))
121 def simulate_vb(n):
122     F_a = 0
123     F_b = 0
124     R = [1]
125     for i in range(n):
126         if vbflip():
127             F_a += 1
128         else:
129             F_b += 1
130         R = np.append(R, ratio(F_a, F_b))
131     return R, F_a
132
133 # simulating 3 times and plotting them
134 plt.figure()
135 for i in range(3):
136     R, F_a = simulate_b(200)

```



```

137     Fa_toss_count_55.append(F_a)
138     plt.plot(n, R, label='$\mathcal{R}$' + str(i+1))
139 plt.xlabel('$F$')
140 plt.ylabel('$R$')
141 plt.title('Trajectory of $\mathcal{R}$ for 3 simulations with $p_a = 0.55$')
142 plt.ylim(-1, 10)
143 plt.legend()
144
145 # for p_a = 0.9
146 plt.figure()
147 for i in range(3):
148     R, F_a = simulate_vb(200)
149     Fa_toss_count_90.append(F_a)
150     plt.plot(n, R, label='$\mathcal{R}$' + str(i+1))
151 plt.xlabel('$F$')
152 plt.ylabel('$R$')
153 plt.title('Trajectory of $\mathcal{R}$ for 3 simulations with $p_a = 0.9$')
154 plt.legend()
155 plt.ylim(0, 1e9)
156 plt.show()
157
158 # %%
159 # posterior distribution for p_a = 0.55
160 def normalconst(F_a):
161     return gm.factorial(F_a) * gm.factorial(200 - F_a) / gm.factorial(201)
162
163 def posterior(F_a, p_a):
164     return p_a ** (F_a) * (1 - p_a) ** (200 - F_a) / normalconst(F_a)
165
166 for i, F_a in enumerate(Fa_toss_count_55):
167     p_a = np.linspace(0, 1, 1000)
168     plt.plot(p_a, posterior(F_a, p_a), label='Posterior' + str(i+1))
169 plt.xlabel('$p_a$')
170 plt.ylabel('Posterior')
171 plt.legend()
172 plt.show()
173
174 # posterior distribution for p_a = 0.9
175 for i, F_a in enumerate(Fa_toss_count_90):
176     p_a = np.linspace(0, 1, 1000)
177     plt.plot(p_a, posterior(F_a, p_a), label='Posterior' + str(i+1))
178 plt.xlabel('$p_a$')
179 plt.ylabel('Posterior')
180 plt.legend()
181 plt.show()
182
183 # %%
184 # back-envelope calculation for p_a close to 0.5
185 # for a Gaussian the width of the error bars is sigma / sqrt(n)
186 # the width should be less than (p_a - 0.5)
187 # p_a - 0.5 > sigma / sqrt(n)
188 # n > sigma ** 2 / (p_a - 0.5) ** 2
189 def backenvelope(sigma):
190     return sigma ** 2 / (0.55 - 0.5) ** 2
191 print(backenvelope(3))
192 print(backenvelope(2))
193
194 l = np.arange(0, 3601)
195 plt.figure()
196 R, F_a = simulate_b(3600)
197 plt.plot(l, R, label='$\mathcal{R}$' + str(i+1))
198 plt.xlabel('$F$')
199 plt.ylabel('$R$')
200 plt.ylim(0, 1e3)
201 plt.legend()
202
203 plt.figure()
204 R, F_a = simulate_b(1600)
205 plt.plot(np.arange(0,1601), R, label='$\mathcal{R}$' + str(i+1))
206 plt.xlabel('$F$')

```

