

IEEE-754

Junseo Shin

April 11, 2024

Introduction

IEEE-754: Floating Point Arithmetic

Introduction

Should We Trust Computers?

- The error in floating point is a form of “noise” in the data.
- The computer automatically “preprocesses” the data in a way that is not always what you want.

How does a computer operate on numbers?

178956970 – 178957034

How does a computer operate on numbers?

178956970 – 178957034



How does a computer operate on numbers?

178956970 – 178957034



= -64

Adding/Subtracting two 32 bit numbers

$$\begin{array}{r} (0)000\ 1010\ 1010\ 1010\ 1010\ 1010\ 1010\ 1010 \\ - (0)000\ 1010\ 1010\ 1010\ 1010\ 1010\ 1110\ 1010 \\ \hline (1)000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0100\ 0000 \end{array}$$

Adding/Subtracting two 32 bit numbers

$$\begin{array}{r} (0)000\ 1010\ 1010\ 1010\ 1010\ 1010\ 1010\ 1010 \\ - (0)000\ 1010\ 1010\ 1010\ 1010\ 1010\ 1110\ 1010 \\ \hline (1)000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0100\ 0000 \end{array}$$

$$= -2^6 = -64$$

Multiplying two binary numbers

$$170 \times 170 = 28900$$

Multiplying two binary numbers

$$170 \times 170 = 28900$$

								1	0	1	0	1	0	1	0
							×	1	0	1	0	1	0	1	0
								0	0	0	0	0	0	0	0
						1	0	1	0	1	0	1	0		
			0	0	0	0	0	0	0	0	0	0			
		1	0	1	0	1	0	1	0	1	0				
	0	0	0	0	0	0	0	0	0	0	0				
	1	0	1	0	1	0	1	0	1	0					
0	0	0	0	0	0	0	0	0	0						
1	0	1	0	1	0	1	0								
1	1	1	0	0	0	0	0	1	1	1	0	0	1	0	0

Dividing two binary numbers

$$\begin{array}{r} 10101010 \\ 10101010 \overline{) 111000011100100} \\ \underline{10101010} \\ 0011011111 \\ \underline{10101010} \\ 0011010100 \\ \underline{10101010} \\ 0010101010 \\ \underline{10101010} \\ 00000000 \end{array}$$

Problems with Binary Integers

No way to represent rational numbers

Problems with Binary Integers

No way to represent rational numbers

- Division is not accurate

$$3/2 = 1.5 \rightarrow 101/10 = 1.\overset{\circ}{1} = 1$$

Problems with Binary Integers

No way to represent rational numbers

- Division is not accurate

$$3/2 = 1.5 \rightarrow 101/10 = 1.\overline{1} \neq 1$$

- Real-Valued Functions (e.g. $\sin(x)$, e^x , \sqrt{x})

IEEE-754: Floating Point Arithmetic

From Decimal to Floating Point

Scientific Notation: $n \times 10^m$

- $1234 = 1.234 \times 10^3$
- $0.01234 = 1.234 \times 10^{-2}$

From Decimal to Floating Point

Scientific Notation: $n \times 10^m$

- $1234 = 1.234 \times 10^3$
- $0.01234 = 1.234 \times 10^{-2}$

Floating Point: (sign, significand, exponent)

$$(-1)^s \times m \times 2^e$$

Converting Decimal to Floating Point

Decimal \rightarrow Binary \rightarrow Floating Point

$$7.45 \rightarrow 7 + 0.45$$

$$\rightarrow 111.01110011001100110011001100110011 \dots$$

1 bit for sign, 8 bits for exponent, 23 bits for significand

Converting Decimal to Floating Point

Decimal \rightarrow Binary \rightarrow Floating Point

$$7.45 \rightarrow 7 + 0.45$$

$$\rightarrow 111.01110011001100110011001100110011 \dots$$

$$\begin{array}{c} \text{sign} \\ \boxed{+} \end{array} \underbrace{1.1101 \dots}_{\text{significand}} \times 2^{\underbrace{2}_{\text{exponent}}}$$

1 bit for sign, 8 bits for exponent, 23 bits for significand

0	1000 0001	1101 1100 1100 1100 1100 110	0110 ...
---	-----------	------------------------------	----------

Machine Epsilon

$$\epsilon = 2^{-(p-1)}$$

- Single Precision: $p = 24$: $\epsilon = 2^{-23} \approx 1.19 \times 10^{-7}$

$$1 + \epsilon = 1.0000\ 0000\ 0000\ 0000\ 0000\ 001$$

- Double Precision: $p = 53$: $\epsilon = 2^{-52} \approx 2.22 \times 10^{-16}$

```
python3
```

```
>>> 0.1 + 0.3
```

```
0.30000000000000004
```

```
>>> 0.3 / 0.10
```

```
2.9999999999999996
```

GNU Multiple Precision Arithmetic Library (GNU MPFR)

- Used in Mathematica
- Follows IEEE-754 standard but with arbitrary precision
- Does this solve the problem?

Stirling's Approximation vs Factorial

$$n! \approx \sqrt{2\pi n} \left(\frac{n}{e}\right)^n$$

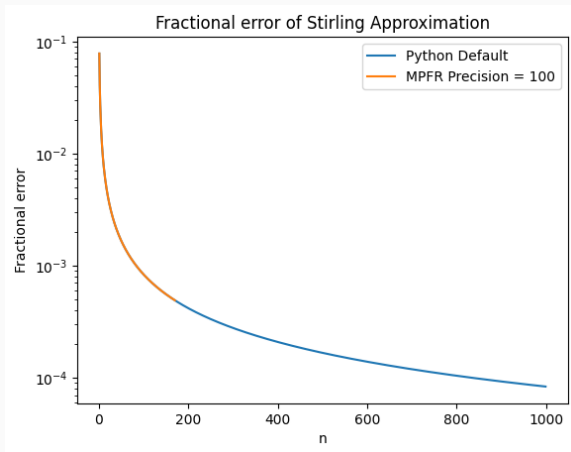
Stirling's Approximation vs Factorial

$$n! \approx \sqrt{2\pi n} \left(\frac{n}{e}\right)^n$$

Python Default (Float64, $p = 53$) vs MPFR ($p = 100$)

- $n = 171 \rightarrow$ Overflow in Float64
- At $n = 100$, Fractional Error of 0.000833 for both!
- Fractional Error of 10^{-6} at $n = 83334$

Float64 vs MPFR



Patriot Defense Missile Failure

- February 25, 1991: Patriot missile defense system failed to intercept SCUD missile
- 28 soldiers died

Simulating Precision Loss

- Integer to 24 bit floating point

Hours	Seconds	Calculated Time (Seconds)	Inaccuracy (Seconds)	Approximate Shift In Range Gate (Meters)
0	0	0	0	0
1	3600	3599.9966	.0034	7
8	28800	28799.9725	.0275	55
20 ^a	72000	71999.9313	.0687	137
48	172800	172799.8352	.1648	330
72	259200	259199.7528	.2472	494
100 ^b	360000	359999.6667	.3433	687

^aContinuous operation exceeding about 20 hours—target outside range gate

^bAlpha Battery ran continuously for about 100 hours

Reference



IEEE standard for floating-point arithmetic.

IEEE Std 754-2019 (Revision of IEEE 754-2008), pages 1–84, 2019.



David Goldberg.

**What every computer scientist should know about
floating-point arithmetic.**

ACM Computing Surveys, 23:5–48, Mar 1991.

[2] [1]

- <https://floating-point-gui.de/>
- [Patriot Missile Defense Software Problem Led to System Failure at Dhahran, Saudi Arabia](#)
- [Floating Point Converter](#)