

DeviceQuery

```
[junseo@r40a-09.sif bin]$ ./DeviceQuery_Template
There is 1 device supporting CUDA
Device 0 name: NVIDIA RTX 4000 Ada Generation
Computational Capabilities: 8.9
Maximum global memory size: 21022244864
Maximum constant memory size: 65536
Maximum shared memory size per block: 49152
Maximum block dimensions: 1024 x 1024 x 64
Maximum grid dimensions: 2147483647 x 65535 x 65535
Warp size: 32
[TIME] [GPU] [Getting GPU Data.] [/home/warehouse/junseo/cuda-code-repo-joonsuuh/Module2/DeviceQuery/template.cu: 14
-55] Elapsed time: 1.5738 ms
[junseo@r40a-09.sif bin]$
```

Figure 0.1: DeviceQuery_template output

Questions

1. What is the compute capability of the NVIDIA Ada architecture?

Compute Capability: 8.9

2. What are the maximum block dimensions for GPUs with the compute capability of this architecture?

Max block dimensions: $1024 \times 1024 \times 64$

3. Suppose you are launching a one dimensional grid and block. If the hardware's maximum grid dimension is 65535 and the maximum block dimension is 1024, what is the maximum number threads can be launched on the GPU?

Max number of threads: $65535 \times 1024 = 67107840$

4. Under what conditions might a programmer choose not want to launch the maximum number of threads?

If multiple kernels are launched at the same time, then a programmer shouldn't launch the maximum number of threads on each kernel function.

5. What can limit a program from launching the maximum number of threads on a GPU?

The memory capacity of the GPU can limit a program from launching the maximum number of threads—e.g, if

6. What is shared memory?

The memory allocated to thread blocks that can be accessed by all threads in the block which allows for low latency access to the data per SM (Streaming Multiprocessor).

7. What is global memory?

Memory that is visible to all thread blocks and specified by the DRAM of the device (GPU)—e.g. a GeForce RTX 4090 has 24 GB of GDDR6X global memory—allocated using `cudaMalloc()` or `cudaMallocManaged()` for CUDA API.

8. What is constant memory?

Similar to global memory i.e. visible to all thread blocks, but is read-only and can't be changed by the threads on the device kernel function.

9. What does warp size signify on a GPU?

It is the number of threads in a block that are assigned/partitioned to an SM. For example, the warp size of the NVIDIA Ada architecture is 32 threads. E.g.

- A one-dimensional block with 64 threads will be partitioned into 2 warps.