

# 1 Module 2

## Vector Add

```
[junseo@r40a-09.sif bin]$ ./VectorAdd_Solution -e ../VectorAdd/Dataset/0/output.raw -i ../VectorAdd/Dataset/0/input0.raw ../VectorAdd/Dataset/0/in
put1.raw -o result0.raw -t vector
[TIME][Generic][Importing data and creating memory on host][/home/warehouse/junseo/cuda-code-repo-joonsuuh/Module3/VectorAdd/solution.cu: 24-30] E
lapsed time: 11.7298 ms
The input length is 16
[TIME][GPU][Allocating GPU memory.][/home/warehouse/junseo/cuda-code-repo-joonsuuh/Module3/VectorAdd/solution.cu: 34-41] Elapsed time: 0.230177 ms
[TIME][GPU][Copying input memory to the GPU.][/home/warehouse/junseo/cuda-code-repo-joonsuuh/Module3/VectorAdd/solution.cu: 43-49] Elapsed time: 0
.04006 ms
[TIME][Compute][Performing CUDA computation][/home/warehouse/junseo/cuda-code-repo-joonsuuh/Module3/VectorAdd/solution.cu: 56-66] Elapsed time: 0.
544428 ms
[TIME][Copy][Copying output memory to the CPU][/home/warehouse/junseo/cuda-code-repo-joonsuuh/Module3/VectorAdd/solution.cu: 68-72] Elapsed time:
0.029234 ms
[TIME][GPU][Freeing GPU Memory][/home/warehouse/junseo/cuda-code-repo-joonsuuh/Module3/VectorAdd/solution.cu: 74-80] Elapsed time: 0.145005 ms
The solution is correct
[junseo@r40a-09.sif bin]$ ./VectorAdd_Solution -e ../VectorAdd/Dataset/0/output.raw -i ../VectorAdd/Dataset/0/input0.raw ../VectorAdd/Dataset/0/in
put1.raw -o result0.raw -t vector
[TIME][Generic][Importing data and creating memory on host][/home/warehouse/junseo/cuda-code-repo-joonsuuh/Module3/VectorAdd/solution.cu: 24-30] E
lapsed time: 11.441 ms
The input length is 16
[TIME][GPU][Allocating GPU memory.][/home/warehouse/junseo/cuda-code-repo-joonsuuh/Module3/VectorAdd/solution.cu: 34-41] Elapsed time: 0.376009 ms
[TIME][GPU][Copying input memory to the GPU.][/home/warehouse/junseo/cuda-code-repo-joonsuuh/Module3/VectorAdd/solution.cu: 43-48] Elapsed time: 0
.03449 ms
[TIME][Compute][Performing CUDA computation][/home/warehouse/junseo/cuda-code-repo-joonsuuh/Module3/VectorAdd/solution.cu: 54-64] Elapsed time: 0.
453247 ms
[TIME][Copy][Copying output memory to the CPU][/home/warehouse/junseo/cuda-code-repo-joonsuuh/Module3/VectorAdd/solution.cu: 66-70] Elapsed time:
0.072061 ms
[TIME][GPU][Freeing GPU Memory][/home/warehouse/junseo/cuda-code-repo-joonsuuh/Module3/VectorAdd/solution.cu: 72-78] Elapsed time: 0.148419 ms
The solution is correct
```

Figure 1.1: VectorAdd\_Solution output

## Questions

1. How many floating operations are being performed in your vector add kernel? EXPLAIN.
2. How many global memory reads are being performed by your kernel? EXPLAIN.
3. How many global memory writes are being performed by your kernel? EXPLAIN.
4. Describe what possible optimizations can be implemented to your kernel to achieve a performance speedup.
5. Name three applications of vector addition.