

Convolution

Questions

- Three applications of convolution:
 - Image processing: e.g. Filters like blurring which can also be used for edge detection via Difference of Gaussians.
 - Audio processing: noise reduction, and EQ filtering.
 - Real-time rendering: e.g. Irradiance environment maps ([GPU Gems 2 Ch 10](#)) which precomputes diffuse reflections from a scene (real or virtual).
- One thread does one multiplication and addition for an element on the mask filter ($2 \times M \times M$), and we have to do this for the color channels $C = 3$ ($C \times 2M^2$). So for an image of width and height $W \times H$, we have $W \times H \times C \times (2M^2)$ float operations.
- The shared memory requires $W \times H \times C$ global reads, but there is an additional of up to

$$4 \times \left(\frac{M-1}{2}\right)^2 + 4 \times \frac{M-1}{2} \times T = (M-1)^2 + 2(M-1) \times T \\ = (M-1) \times (M-1 + 2T)$$

global reads per tile (where T is the tile size $T \times T$) because if a halo cell is not a ghost cell, it has to be read from global memory from the else statement in the nested for loop.

- The global writes is equal to the number of output pixels we need to write to, so $W \times H \times C$ global writes.
- Minimum: At the corner pixel we would have only 9 out of the 25 mask filter elements performing real operation. So we have $9 \times 2 = 18$ additions and multiplications and $18 \times 3 = 54$ float operations for the three color channels.

Maximum: At the center pixel we would have all mask elements performing real operations, so $25 \times 2 \times 3 = 150$ float operations.

Average: We have alot of pixels to deal with...

- 4 corner pixels performing 54float operations
- 2×4 pixels horizontally and vertically adjacent to the corner pixels performing $(3 \times 4) \times 6 = 72$ float operations
- $4 \times (N-4)$ edge pixels performing $(3 \times 5) \times 6 = 90$ float operations
- 4 pixels diagonally adjacent to the corner pixels performing $(4 \times 4) \times 6 = 96$ float operations
- $4 \times (N-4)$ inner edge pixels performing $(4 \times 5) \times 6 = 120$ float operations.
- $(N-4)^2$ inner pixels performing $(5 \times 5) \times 6 = 150$ float operations.

Double checking if we counted the number of pixels correctly:

$$4 + 2 \times 4 + 4 \times (N-4) + 4 + 4 \times (N-4) + (N-4)^2 = 16 + 8N - 32 + N^2 - 8N + 16 = N^2$$

So the average float operations per pixel is

$$\frac{4 \times 54 + 8 \times 72 + 4 \times (N-4) \times 90 + 4 \times 96 + 4 \times (N-4) \times 120 + (N-4)^2 \times 150}{N^2}$$

- For a $W \times H = 64 \times 64$ with a computation time of 0.066 622 ms, the throughput is roughly

$$\frac{64 \times 64 \times 3 \times 50 \text{ FLO}}{0.066 \text{ 622 ms}} = 9.2 \text{ GFLOPS}$$

This would scale linearly with the size of the input image or $O(W \times H)$

7. The asymptotic time complexity of the sequential convolution is still $O(W \times H)$ because the parallel convolution is only changing the constant factor by adding parallel execution threads (P) that will roughly decrease the computation time by the $1/P$.
8. For the 64 input image, it take 0.242 ms to allocate the memory on the GPU, 0.402 ms to copy the input image to the GPU, and 0.058 ms to copy the output image back to the CPU. This overhead will scale linearly with the output size, since the host-device memory transfer have a limit to the bandwidth given by the PCIe bus.
9. If the mask size was increased to $M = 1024$ with the block dimension 16×16 , the shared memory would be ineffective because most of the time would be spend doing global reads. Futhermore, the constant memory capacity of 64KB would be exceeded ($1024^2 \times 4 \approx 4 \text{ MB}$), so there would be a CUDA runtime error or compilation error that would prevent the kernel from running.
10. There must be a separate output memory buffer, because the convolution needs to access the original input values, so doing the convolution in place would overwrite the original input values and threads would try to read modified values and run into race conditions.
11. An identity mask does not change the input image once the mask is applied through convolution, and it has the value of 1 at the center and 0 elsewhere. So for the $M = 5$ mask, the identity mask is

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

OUTPUT

[illegible]

Figure 1: Convolution