

Reduction

Questions

1. 3 Applications of reduction:

- (a) Prefix sum: SAT (Summed Area Table) which can be used for depth of field blurring.
- (b) Computing single values such as max/min, average, or standard deviation.
- (c) Physics simulations: e.g. N-body simulations where we need to compute the sum of forces acting on a particle.

2. Since the loading into shared memory does one addition per thread, and in the reduction step we have $n = \log_2 \text{blockDim.x}$ steps doing one addition times the stride or as a finite geometric sum

$$\begin{aligned} \frac{\text{blockDim.x}}{2} + \frac{\text{blockDim.x}}{4} + \dots + 1 &= a \frac{1 - r^n}{1 - r} \\ &= \frac{\text{blockDim.x}}{2} \frac{1 - (1/2)^{\log_2 \text{blockDim.x}}}{1 - 1/2} \\ &= \text{blockDim.x} \left(1 - \frac{1}{\text{blockDim.x}} \right) \\ &= \text{blockDim.x} - 1 \end{aligned}$$

so we have $(2 \times \text{blockDim.x} - 1) \times \text{gridDim.x}$ additions in total.

- 3. All the reads are done in the shared memory step which loads 2 elements per thread, so we have $2 \times \text{blockDim.x} \times \text{gridDim.x}$ reads in total or $2 \times N$ for an input size N .
- 4. The writes are done once per block when the `threadIdx.x == 0` so we have gridDim.x (number of blocks) writes in total.
- 5. Min: If `i >= len` then we would have 0 real operations. Max: The thread would have 1 addition in the shared step and $\log_2 \text{blockDim.x}$ additions in the reduction step. Average: The average number of operations is the result in a block size from 2. divided by the number of thread in a block:

$$\frac{2 \times \text{blockDim.x} - 1}{\text{blockDim.x}} = 2 - \frac{1}{\text{blockDim.x}}$$

- 6. The thread block will synch $\lfloor \log_2 \text{blockDim.x} \rfloor$ times in the reduction step.
- 7. The shared memory reduces the the number of global memory accesses to zero in the reduction step!
- 8. We could use segmented multiblock sums and/or thread coarsening for faster DRAM access.
- 9. We could modify the kernel with a grid-stride loop to really large input sizes that are bigger than the max dimensions, or use the segmented multiblock sum.
- 10. No, we can not use non-commutative operation because the threads will be mixing up the order of the operations as it runs in parallel.
- 11. Yes, the floating point arithmetic is not associative, so the order of the operations will affect the result.

OUPUT

```
[junseo@r40a-09.sif Reduction]$ ./reduction.sh
0 The solution is correct
1 The solution is correct
2 The solution is correct
3 The solution is correct
4 The solution is correct
5 The solution is correct
6 The solution is correct
7 The solution is correct
8 The solution is correct
9 The solution is correct
10 The solution is correct
11 The solution is correct
```

Figure 1: Successes